

Operating System Project #5: File System

Global Economics / Computer Science
2014311577 Dongmin Kim

Fixing Bmap Function

1. Fixing struct types

First of all, to support doubly-indirect block for the file system, we need to fix the inode structures for both on-disk and on memory.

On-disk inode is defined in fs.h, as follows.

```
#define NDIRECT 11
#define NINDIRECT (BSIZE / sizeof(uint))
#define NDINDIRECT (BSIZE / sizeof(uint)) * (BSIZE / sizeof(uint))
#define MAXFILE (NDIRECT + NINDIRECT + NDINDIRECT)

// On-disk inode structure
struct dinode {
    short type;           // File type
    short major;          // Major device number (T_DEV only)
    short minor;          // Minor device number (T_DEV only)
    short nlink;           // Number of links to inode in file system
    uint size;             // Size of file (bytes)
    uint addrs[NDIRECT+2]; // Data block addresses
};
```

Here, number of direct block(NDIRECT) is 11, and indirect block(NINDIRECT) is (BSIZE / sizeof(uint)), which is 128 (512/4). Similarly, it is possible to define doubly-indirect block. Since doubly indirect block contains pointer to the block of pointers, so it can cover up to (BSIZE/sizeof(uint)) * (BSIZE/sizeof(uint)), which is 16384. So in total, this system can cover up to 16523 sectors.

So the difference is the structure of diode->addrs, since it allows 11 direct/ 1 indirect/ 1 double-indirect pointers. Therefore the size is NDIRECT + 2.

similarly, in-memory copy can be fixed as follows, and this is defined in file.h

```
// in-memory copy of an inode
struct inode {
    uint dev;           // Device number
    uint inum;          // Inode number
    int ref;            // Reference count
    struct sleeplock lock; // protects everything below here
    int valid;          // inode has been read from disk?

    short type;         // copy of disk inode
    short major;
    short minor;
```

```

short nlink;
uint size;
uint addrs[NDIRECT+2];
};

```

2. bmap function

bmap function is modified as follows:

```

static uint
bmap(struct inode *ip, uint bn)
{
    uint addr, *a, *a2;
    struct buf *bp, *bp2;
    uint idx1, idx2;

    if(bn < NDIRECT){
        if((addr = ip->addrs[bn]) == 0)
            ip->addrs[bn] = addr = balloc(ip->dev);
        return addr;
    }
    bn -= NDIRECT;

    if(bn < NINDIRECT){
        // Load indirect block, allocating if necessary.
        if((addr = ip->addrs[NDIRECT]) == 0)
            ip->addrs[NDIRECT] = addr = balloc(ip->dev);
        bp = bread(ip->dev, addr);
        a = (uint*)bp->data;
        if((addr = a[bn]) == 0){
            a[bn] = addr = balloc(ip->dev);
            log_write(bp);
        }
        brelse(bp);
        return addr;
    }

    bn -= NINDIRECT;

    idx1 = (bn) / (BSIZE/sizeof(uint));
    idx2 = (bn) % (BSIZE/sizeof(uint));

    if(bn < NDINDIRECT){
        // Load double indirect block, allocating if necessary.

        if((addr = ip->addrs[NDIRECT+1]) == 0)
            ip->addrs[NDIRECT+1] = addr = balloc(ip->dev);
    }
}

```

```

bp = bread(ip->dev, addr);
a = (uint*)bp->data;

if((addr = a[idx1]) == 0){
    a[idx1] = addr = balloc(ip->dev);
    log_write(bp);
}
brelse(bp);

bp2 = bread(ip->dev, addr);
a2 = (uint*)bp2->data;

if((addr = a2[idx2]) == 0){
    a2[idx2] = addr = balloc(ip->dev);
    log_write(bp2);
}

brelse(bp2);
return addr;
}

panic("bmap: out of range");
}

```

First two parts allocates Direct and Indirect blocks.

```

if(bn < NDIRECT){
    if((addr = ip->addrs[bn]) == 0)
        ip->addrs[bn] = addr = balloc(ip->dev);
    return addr;
}
bn -= NDIRECT;

if(bn < NINDIRECT){
    // Load indirect block, allocating if necessary.
    if((addr = ip->addrs[NDIRECT]) == 0)
        ip->addrs[NDIRECT] = addr = balloc(ip->dev);
    bp = bread(ip->dev, addr);
    a = (uint*)bp->data;
    if((addr = a[bn]) == 0){
        a[bn] = addr = balloc(ip->dev);
        log_write(bp);
    }
    brelse(bp);
    return addr;
}

```

Similar to those processes above, allocation for doubly indirect block can be done as follows:

```
bn -= NINDIRECT;

idx1 = (bn) / (BSIZE/sizeof(uint));
idx2 = (bn) % (BSIZE/sizeof(uint));

if(bn < NDINDIRECT){
    // Load double indirect block, allocating if necessary.

    if((addr = ip->addrs[NINDIRECT+1]) == 0)
        ip->addrs[NINDIRECT+1] = addr = balloc(ip->dev);

    bp = bread(ip->dev, addr);
    a = (uint*)bp->data;

    if((addr = a[idx1]) == 0){
        a[idx1] = addr = balloc(ip->dev);
        log_write(bp);
    }
    brelse(bp);

    bp2 = bread(ip->dev, addr);
    a2 = (uint*)bp2->data;

    if((addr = a2[idx2]) == 0){
        a2[idx2] = addr = balloc(ip->dev);
        log_write(bp2);
    }

    brelse(bp2);
    return addr;
}

panic("bmap: out of range");
```

Firstly, last index of the pointer array in dinode is referred, and allocate a block if necessary. And two reference numbers, which are idx1, idx2 indicates their own location in first/ and second pointer blocks. Therefore, by calling a[idx1], it is possible to locate and allocate the blocks that contains pointer blocks. Similarly, by calling a2[idx2], it is possible to locate and allocate the blocks that contains actual blocks.

3. Testing

For the testing, FSSIZE were set to 21113.

```
dongmin@dongmin-ThinkPad-E495:~/Projects/2020_Spring/OS_2020_1st_swe3004/xv6 projects/[OS project-5] 2014311577_김동민_ok$ make qemu-nox
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 2 -m 512
xv6...
cpu1: starting 1
8010a460 2c
cpu0: starting 0
sb: size 21113 nblocks 21049 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
Student ID: 2014311577
Name: Dongmin Kim
Message from developer: Welcome to xv6 OS!
$ test
.....
.....break
wrote 16523 sectors
done; ok
$
```